

```
In [21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import (confusion_matrix, classification_report,
                             roc_auc_score, roc_curve, accuracy_score)
import shap
import joblib
```

```
In [ ]: #Load the dataset
df = pd.read_csv('blood.csv')
print("Rows, Columns:", df.shape)
df.head()
```

Rows, Columns: (748, 5)

	Recency	Frequency	Monetary	Time	Class
0	2	50	12500	99	1
1	0	13	3250	28	1
2	1	17	4000	36	1
3	2	20	5000	45	1
4	1	24	6000	77	0

```
In [ ]: #Quick exploration
print("Columns and types:")
print(df.dtypes)
print("\nMissing values per column:")
print(df.isnull().sum())
print("\nBasic stats for numeric columns:")
display(df.describe().T)
```

```
Columns and types:
Recency      int64
Frequency    int64
Monetary     int64
Time         int64
Class        int64
dtype: object
```

```
Missing values per column:
Recency      0
Frequency    0
Monetary     0
Time         0
Class        0
dtype: int64
```

Basic stats for numeric columns:

	count	mean	std	min	25%	50%	75%	max
Recency	748.0	9.506684	8.095396	0.0	2.75	7.0	14.0	74.0
Frequency	748.0	5.516043	5.841825	1.0	2.00	4.0	7.0	50.0
Monetary	748.0	1378.676471	1459.826781	250.0	500.00	1000.0	1750.0	12500.0
Time	748.0	34.284759	24.380307	2.0	16.00	28.0	50.0	99.0
Class	748.0	0.237968	0.426124	0.0	0.00	0.0	0.0	1.0

```
In [25]: if 'recommendation' in df.columns:
    target_col = 'recommendation'
elif 'disease' in df.columns:
    target_col = 'disease'
elif 'diabetes' in df.columns:
    target_col = 'diabetes'
else:
    numeric_cols = df.select_dtypes(include=np.number).columns.tolist()
    # heuristics - adapt to dataset columns
    risk = pd.Series(0, index=df.index)
    if 'cholesterol' in df.columns:
        risk = risk | (df['cholesterol'] > 240)
    if 'blood_glucose' in df.columns:
        risk = risk | (df['blood_glucose'] > 126)
    for c in ['systolic_bp', 'blood_pressure', 'bp_sys']:
        if c in df.columns:
            risk = risk | (df[c] > 140)
    df['risk_label'] = risk.astype(int)
    target_col = 'risk_label'
print("Created synthetic target 'risk_label' - please adapt if you have a real
```

Created synthetic target 'risk_label' - please adapt if you have a real label.



```
In [26]: #EDA
#Missing value heatmap
plt.figure(figsize=(10,3))
sns.heatmap(df.isnull(), cbar=False)
```

```


plt.title("Missing values (black = present, yellow=missing)")
plt.show()

#Numeric distributions for top numeric columns
num_cols = df.select_dtypes(include=np.number).columns.tolist()
num_cols = [c for c in num_cols if c != target_col][:6] # show up to 6
for c in num_cols:
    plt.figure()
    sns.histplot(df[c].dropna(), kde=True)
    plt.title(f"Distribution: {c}")
    plt.xlabel(c)
    plt.ylabel("count")
    plt.show()


```




Distribution: Frequency




Distribution: Monetary





```
In [27]: #Correlation heatmap
corr = df.select_dtypes(include=np.number).corr()
plt.figure(figsize=(10,8))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="RdYlBu_r", square=False)
plt.title("Correlation matrix (numeric features)")
plt.show()
```



```
In [ ]: #Preprocessing pipeline
feature_cols = [c for c in df.columns if c != target_col]

numeric_features = df[feature_cols].select_dtypes(include=np.number).columns.tolist
categorical_features = [c for c in feature_cols if c not in numeric_features]

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ],
    remainder='drop' # drop other columns if any
)
```

```
In [55]: #Train/test split
X = df[feature_cols]
y = df[target_col]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
print("Train/test sizes:", X_train.shape, X_test.shape)
```

Train/test sizes: (598, 5) (150, 5)

```
In [ ]: #Model training: RandomForest , XGBoost
y_train = y_train.astype(int)
y_test = y_test.astype(int)

rf_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(n_estimators=200, random_state=42))
])

xgb_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', XGBClassifier(
        objective='binary:logistic',
        eval_metric='logloss',
        use_label_encoder=False,
        base_score=0.5,
        random_state=42
    ))
])
```

```
In [ ]: # Cross-validate quickly
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

print("RF CV accuracy:",
      cross_val_score(rf_pipeline, X_train, y_train, cv=cv, scoring='accuracy').mean())

print("XGB CV accuracy:",
      cross_val_score(xgb_pipeline, X_train, y_train, cv=cv, scoring='accuracy').mean())

# Fit best model
best_pipeline = rf_pipeline
best_pipeline.fit(X_train, y_train)
```

RF CV accuracy: 1.0

XGB CV accuracy: 1.0

Out[]:

Pipeline

▼ Parameters

steps	[(<code>'preprocessor'</code> , ...), (<code>'clf'</code> , ...)]
transform_input	None
memory	None
verbose	False

preprocessor: ColumnTransformer

▼ Parameters


transformers	[('num', ...), ('cat', ...)]
remainder	'drop'
sparse_threshold	0.3
n_jobs	None
transformer_weights	None
verbose	False
verbose_feature_names_out	True
force_int_remainder_cols	'deprecated'

▼ num cat

`['Recency', 'Frequency', 'Monetary', 'Time', 'Class']`

SimpleImputer	SimpleImpute
▼ Parameters	▼ Parameters
missing_values	nan
strategy	'median'
fill_value	None
copy	True
add_indicator	False
keep_empty_features	False

StandardScaler	OneHotEncode
-----------------------	---------------------



In []:

```
y_proba = None
try:
    y_proba = best_pipeline.predict_proba(X_test)[:,1]
except Exception:
    pass

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification report:")
print(classification_report(y_test, y_pred))
```


Accuracy: 1.0

Classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	150
accuracy			1.00	150
macro avg	1.00	1.00	1.00	150
weighted avg	1.00	1.00	1.00	150

In []:

```
# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [ ]: #ROC Curve
best_pipeline.fit(X_train, y_train)

if hasattr(best_pipeline, "predict_proba") and len(np.unique(y_test)) == 2:
    y_proba = best_pipeline.predict_proba(X_test)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, y_proba)
    auc = roc_auc_score(y_test, y_proba)
    plt.figure(figsize=(6,5))
    plt.plot(fpr, tpr, label=f"AUC={auc:.3f}", linewidth=2)
    plt.plot([0,1],[0,1], '--', alpha=0.6)
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title("ROC Curve")
    plt.legend()
    plt.grid(alpha=0.3)
    plt.show()
else:
    if not hasattr(best_pipeline, "predict_proba"):
        print("Model does not support probability predictions.")
    else:
        print("Cannot plot ROC - y_test does not contain both classes:", np.unique(y_test))
```

Cannot plot ROC - y_test does not contain both classes: [0]

```
In [ ]: #Feature importance
pre = best_pipeline.named_steps['preprocessor']

num_feats = numeric_features

cat_feats = []
if categorical_features:
    ohe = pre.named_transformers_['cat'].named_steps['onehot']
```


```

        cat_names = ohe.get_feature_names_out(categorical_features)
        cat_feats = list(cat_names)
all_feats = num_feats + cat_feats

clf = best_pipeline.named_steps['clf']
importances = clf.feature_importances_
feat_imp = pd.Series(importances, index=all_feats).sort_values(ascending=False).head(20)

plt.figure(figsize=(8,6))
sns.barplot(x=feat_imp.values, y=feat_imp.index)
plt.title("Top 20 Feature Importances")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()


```



```
In [ ]: #SHAP explanations
explainer = shap.Explainer(clf, pre.transform(X_train).astype(float))

X_shap = pre.transform(X_test).astype(float)
shap_values = explainer(X_shap[:200])

shap.summary_plot(shap_values, features=X_shap[:200], feature_names=all_feats, show
```



```
In [ ]: #Recommendation generator
def generate_recommendation(single_row_df):
    """
    Input: single row pandas DataFrame with the same columns as original X
    Output: human readable recommendation
    """
    pred = best_pipeline.predict(single_row_df)[0]

    mapping = {
        0: "Low risk: Maintain current healthy habits; annual checkup recommended."
        1: "Elevated risk: Schedule a doctor visit for assessment; consider lifestyle changes"
    }

    return mapping.get(pred, f"Predicted label {pred}. Consult clinician for personal evaluation.")

example = X_test.iloc[[0]]
print(generate_recommendation(example))
```

Low risk: Maintain current healthy habits; annual checkup recommended.